

# Homework-2

December 6, 2022

## 1 Designing a backdoor detector for BadNets trained on the YouTube Face dataset using the pruning defense.

```
[1]: # All necessary imports
import os
import tarfile
import requests
import re
import sys
import warnings
warnings.filterwarnings('ignore')
import h5py
import numpy as np
import tensorflow as tf
from tensorflow import keras
from keras import backend as K
from keras.models import Model
import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1.inset_locator import inset_axes
import matplotlib.font_manager as font_manager
import cv2
```

Define function to load the data

```
[2]: # Load data
def data_loader(filepath):
    data = h5py.File(filepath, 'r')
    x_data = np.array(data['data'])
    y_data = np.array(data['label'])
    x_data = x_data.transpose((0,2,3,1))
    return x_data, y_data
```

Follow instructions under [Data Section](#) to download the datasets.

We will be using the clean validation data (valid.h5) from cl folder to design the defense and clean test data (test.h5 from cl folder) and sunglasses poisoned test data (bd\_test.h5 from bd folder) to evaluate the models.

```
[3]: ## To-do ##
# After downloading the datasets, provide corresponding filepaths below

clean_data_valid_filename = "./data/cl/valid.h5"

clean_data_test_filename = "./data/cl/test.h5"
poisoned_data_test_filename = "./data/bd/bd_test.h5"
```

Read the data:

```
[4]: cl_x_valid, cl_y_valid = data_loader(clean_data_valid_filename)

cl_x_test, cl_y_test = data_loader(clean_data_test_filename)
bd_x_test, bd_y_test = data_loader(poisoned_data_test_filename)
```

Visualizing the clean test data

```
[5]: # Plot some images from the validation set (see https://mrdatascience.com/
↪how-to-plot-mnist-digits-using-matplotlib/)
num = 10
np.random.seed(45)
randIdx = [np.random.randint(10000) for i in range(num)]
num_row = 2
num_col = 5# plot images
fig, axes = plt.subplots(num_row, num_col, figsize=(3*num_col,3*num_row))
for i in range(num):
    ax = axes[i//num_col, i%num_col]
    ax.imshow(cl_x_test[randIdx[i]].astype('uint8'))
    ax.set_title('label: {:.0f}'.format(cl_y_test[randIdx[i]]))
    ax.set_xticks([])
    ax.set_yticks([])
plt.tight_layout()
plt.show()
```



Visualizing the sunglasses poisoned test data

```
[6]: # Plot some images from the validation set (see https://mrdatascience.com/
      ↪how-to-plot-mnist-digits-using-matplotlib/)
num = 10
np.random.seed(45)
randIdx = [np.random.randint(10000) for i in range(num)]
num_row = 2
num_col = 5 # plot images
fig, axes = plt.subplots(num_row, num_col, figsize=(3*num_col,3*num_row))
for i in range(num):
    ax = axes[i//num_col, i%num_col]
    ax.imshow(bd_x_test[randIdx[i]].astype('uint8'))
    ax.set_title('label: {:.0f}'.format(bd_y_test[randIdx[i]]))
    ax.set_xticks([])
    ax.set_yticks([])
plt.tight_layout()
plt.show()
```



Load the backdoored model.

The backdoor model and its weights can be found [here](#)

```
[7]: ## To-do ##

# First create clones of the original badnet model (by providing the model_
      ↪filepath below)
# The result of repairing B_clone will be B_prime

B = keras.models.load_model("./model/bd_net.h5")
B.load_weights("./model/bd_weights.h5")
```

```
B_clone = keras.models.load_model("./model/bd_net.h5")
B_clone.load_weights("./model/bd_weights.h5")
```

Output of the original badnet accuracy on the validation data:

```
[8]: # Get the original badnet model's (B) accuracy on the validation data
cl_label_p = np.argmax(B(cl_x_valid), axis=1)
clean_accuracy = np.mean(np.equal(cl_label_p, cl_y_valid)) * 100

print("Clean validation accuracy before pruning {0:3.6f}".
      ↪format(clean_accuracy))

K.clear_session()
```

Clean validation accuracy before pruning 98.649000

```
[9]: print(B.summary())
```

Model: "model\_1"

```
-----
Layer (type)                 Output Shape          Param #   Connected to
=====
input (InputLayer)           [(None, 55, 47, 3)]   0         []
conv_1 (Conv2D)               (None, 52, 44, 20)    980       ['input[0][0]']
pool_1 (MaxPooling2D)         (None, 26, 22, 20)    0         ['conv_1[0][0]']
conv_2 (Conv2D)               (None, 24, 20, 40)    7240      ['pool_1[0][0]']
pool_2 (MaxPooling2D)         (None, 12, 10, 40)    0         ['conv_2[0][0]']
conv_3 (Conv2D)               (None, 10, 8, 60)     21660     ['pool_2[0][0]']
pool_3 (MaxPooling2D)         (None, 5, 4, 60)      0         ['conv_3[0][0]']
conv_4 (Conv2D)               (None, 4, 3, 80)      19280     ['pool_3[0][0]']
flatten_1 (Flatten)           (None, 1200)          0         []
```

```

['pool_3[0][0]']

flatten_2 (Flatten)          (None, 960)          0
['conv_4[0][0]']

fc_1 (Dense)                 (None, 160)         192160
['flatten_1[0][0]']

fc_2 (Dense)                 (None, 160)         153760
['flatten_2[0][0]']

add_1 (Add)                  (None, 160)          0          ['fc_1[0][0]',
                                                    'fc_2[0][0]']

activation_1 (Activation)    (None, 160)          0          ['add_1[0][0]']

output (Dense)               (None, 1283)        206563
['activation_1[0][0]']

=====
=====
Total params: 601,643
Trainable params: 601,643
Non-trainable params: 0

-----
-----
None

```

Write code to implement pruning defense

```

[10]: ## To-do ##

# Redefine model to output right after the last pooling layer ("pool_3")
intermediate_model = Model(inputs=B.inputs, outputs=B.get_layer('pool_3').
    ↪output)

# Get feature map for last pooling layer ("pool_3") using the clean validation_
↪data and intermediate_model
feature_maps_cl = intermediate_model(cl_x_valid)

# Get average activation value of each channel in last pooling layer ("pool_3")
averageActivationsCl = np.mean(feature_maps_cl, 0)

# Store the indices of average activation values (averageActivationsCl) in_
↪increasing order
avgActByCh = np.mean(np.abs(feature_maps_cl), axis = (0, 1, 2))
idxToPrune = np.argsort(np.abs(avgActByCh))

```

```

# Get the conv_4 layer weights and biases from the original network that will
→ be used for pruning
# Hint: Use the get_weights() method (https://stackoverflow.com/questions/
→ 43715047/how-do-i-get-the-weights-of-a-layer-in-keras)
lastConvLayerWeights = B_clone.layers[7].get_weights()[0]
lastConvLayerBiases = B_clone.layers[7].get_weights()[1]

flag = [0, 0, 0]

for chIdx in idxToPrune:

    # Prune one channel at a time
    # Hint: Replace all values in channel 'chIdx' of lastConvLayerWeights and
    → lastConvLayerBiases with 0
    lastConvLayerWeights[:, :, :, chIdx] = 0
    lastConvLayerBiases[chIdx] = 0

    # Update weights and biases of B_clone
    # Hint: Use the set_weights() method
    B_clone.layers[7].set_weights([lastConvLayerWeights, lastConvLayerBiases])

    # Evaluate the updated model's (B_clone) clean validation accuracy
    cl_label_p_valid = np.argmax(B_clone(cl_x_valid), axis=1)
    clean_accuracy_valid = np.mean(np.equal(cl_label_p_valid, cl_y_valid)) * 100

    # If drop in clean_accuracy_valid is just greater (or equal to) than the
    → desired threshold compared to clean_accuracy, then save B_clone as B_prime
    → and break
    if clean_accuracy - clean_accuracy_valid >= 2 and not flag[0]:
        B_clone.save('./fixed_models/bd_2.h5')
        B_clone.save_weights('./fixed_models/bd_2_weights.h5')
        print("Model has been saved as bd_2.h5 and bd_2_weights.h5")
        flag[0] = 1
    if clean_accuracy - clean_accuracy_valid >= 4 and not flag[1]:
        B_clone.save('./fixed_models/bd_4.h5')
        B_clone.save_weights('./fixed_models/bd_4_weights.h5')
        print("Model has been saved as bd_4.h5 and bd_4_weights.h5")
        flag[1] = 1
    if clean_accuracy - clean_accuracy_valid >= 10 and not flag[2]:
        B_clone.save('./fixed_models/bd_10.h5')
        B_clone.save_weights('./fixed_models/bd_10_weights.h5')
        print("Model has been saved as bd_10.h5 and bd_10_weights.h5")
        flag[2] = 1
        break

```

Model has been saved as bd\_2.h5 and bd\_2\_weights.h5

Now we need to combine the models into a repaired goodnet G that outputs the correct class if the

test input is clean and class N+1 if the input is backdoored. One way to do it is to “subclass” the models in Keras:

```
[11]: #https://stackoverflow.com/questions/64983112/
      ↪keras-vertical-ensemble-model-with-condition-in-between
class G(tf.keras.Model):
    def __init__(self, B, B_prime):
        super(G, self).__init__()
        self.B = B
        self.B_prime = B_prime

    def predict(self, data):
        y = np.argmax(self.B(data), axis=1)
        y_prime = np.argmax(self.B_prime(data), axis=1)
        tmpRes = np.array([y[i] if y[i] == y_prime[i] else 1283 for i in
        ↪range(y.shape[0])])
        res = np.zeros((y.shape[0], 1284))
        res[np.arange(tmpRes.size), tmpRes] = 1
        return res

    # For small amount of inputs that fit in one batch, directly using call()
    ↪is recommended for faster execution,
    # e.g., model(x), or model(x, training=False) is faster than model.
    ↪predict(x) and do not result in
    # memory leaks (see for more details https://www.tensorflow.org/api_docs/
    ↪python/tf/keras/Model#predict)
    def call(self, data):
        y = np.argmax(self.B(data), axis=1)
        y_prime = np.argmax(self.B_prime(data), axis=1)
        tmpRes = np.array([y[i] if y[i] == y_prime[i] else 1283 for i in
        ↪range(y.shape[0])])
        res = np.zeros((y.shape[0], 1284))
        res[np.arange(tmpRes.size), tmpRes] = 1
        return res
```

However, Keras prevents from saving this kind of subclassed model as HDF5 file since it is not serializable. However, we still can use this architecture for model evaluation.

Load the saved B\_prime model

```
[12]: ## To-do ##
      # Provide B_prime model filepath below

      # 2%
      B_prime_2 = keras.models.load_model("./fixed_models/bd_2.h5")
      B_prime_2.load_weights("./fixed_models/bd_2_weights.h5")

      # 4%
```

```

B_prime_4 = keras.models.load_model("./fixed_models/bd_4.h5")
B_prime_4.load_weights("./fixed_models/bd_4_weights.h5")

# 10%
B_prime_10 = keras.models.load_model("./fixed_models/bd_10.h5")
B_prime_10.load_weights("./fixed_models/bd_10_weights.h5")

```

Check performance of the repaired model on the test data:

```

[13]: # 2%
cl_label_p = np.argmax(B_prime_2.predict(cl_x_test), axis=1)
clean_accuracy_B_prime_2 = np.mean(np.equal(cl_label_p, cl_y_test))*100
print('Clean Classification accuracy for B_prime_2:', clean_accuracy_B_prime_2)

bd_label_p = np.argmax(B_prime_2.predict(bd_x_test), axis=1)
asr_B_prime_2 = np.mean(np.equal(bd_label_p, bd_y_test))*100
print('Attack Success Rate for B_prime_2:', asr_B_prime_2)
print("-----")

# 4%
cl_label_p = np.argmax(B_prime_4.predict(cl_x_test), axis=1)
clean_accuracy_B_prime_4 = np.mean(np.equal(cl_label_p, cl_y_test))*100
print('Clean Classification accuracy for B_prime_4:', clean_accuracy_B_prime_4)

bd_label_p = np.argmax(B_prime_4.predict(bd_x_test), axis=1)
asr_B_prime_4 = np.mean(np.equal(bd_label_p, bd_y_test))*100
print('Attack Success Rate for B_prime_4:', asr_B_prime_4)
print("-----")

# 10%
cl_label_p = np.argmax(B_prime_10.predict(cl_x_test), axis=1)
clean_accuracy_B_prime_10 = np.mean(np.equal(cl_label_p, cl_y_test))*100
print('Clean Classification accuracy for B_prime_10:',
      ↪clean_accuracy_B_prime_10)

bd_label_p = np.argmax(B_prime_10.predict(bd_x_test), axis=1)
asr_B_prime_10 = np.mean(np.equal(bd_label_p, bd_y_test))*100
print('Attack Success Rate for B_prime_10:', asr_B_prime_10)

```

```

Clean Classification accuracy for B_prime_2: 96.31332813717849
Attack Success Rate for B_prime_2: 100.0
-----

```

```

Clean Classification accuracy for B_prime_4: 92.29150428682775
Attack Success Rate for B_prime_4: 99.98441153546376
-----

```

```

Clean Classification accuracy for B_prime_10: 84.54403741231489
Attack Success Rate for B_prime_10: 77.20966484801247

```



Check performance of the original model on the test data:

```
[14]: cl_label_p = np.argmax(B.predict(cl_x_test), axis=1)
      clean_accuracy_B = np.mean(np.equal(cl_label_p, cl_y_test))*100
      print('Clean Classification accuracy for B:', clean_accuracy_B)

      bd_label_p = np.argmax(B.predict(bd_x_test), axis=1)
      asr_B = np.mean(np.equal(bd_label_p, bd_y_test))*100
      print('Attack Success Rate for B:', asr_B)
```

Clean Classification accuracy for B: 98.62042088854248

Attack Success Rate for B: 100.0

Create repaired network

```
[15]: # Repaired network repaired_net
      # 2%
      repaired_net_2 = G(B, B_prime_2)

      # 4%
      repaired_net_4 = G(B, B_prime_4)

      # 10%
      repaired_net_10 = G(B, B_prime_10)
```

Check the performance of the repaired\_net on the test data

```
[16]: # 2%
      cl_label_p = np.argmax(repaired_net_2(cl_x_test), axis=1)
      clean_accuracy_repaired_net = np.mean(np.equal(cl_label_p, cl_y_test))*100
      print('Clean Classification accuracy for repaired_net_2:',
            ↪clean_accuracy_repaired_net)

      bd_label_p = np.argmax(repaired_net_2(bd_x_test), axis=1)
      asr_repaired_net = np.mean(np.equal(bd_label_p, bd_y_test))*100
      print('Attack Success Rate for repaired_net_2:', asr_repaired_net)
      print("-----")

      # 4%
      cl_label_p = np.argmax(repaired_net_4(cl_x_test), axis=1)
      clean_accuracy_repaired_net = np.mean(np.equal(cl_label_p, cl_y_test))*100
      print('Clean Classification accuracy for repaired_net_4:',
            ↪clean_accuracy_repaired_net)

      bd_label_p = np.argmax(repaired_net_4(bd_x_test), axis=1)
      asr_repaired_net = np.mean(np.equal(bd_label_p, bd_y_test))*100
      print('Attack Success Rate for repaired_net_4:', asr_repaired_net)
      print("-----")
```

```

# 10%
cl_label_p = np.argmax(repaired_net_10(cl_x_test), axis=1)
clean_accuracy_repaired_net = np.mean(np.equal(cl_label_p, cl_y_test))*100
print('Clean Classification accuracy for repaired_net_10:',
      ↪clean_accuracy_repaired_net)

bd_label_p = np.argmax(repaired_net_10(bd_x_test), axis=1)
asr_repaired_net = np.mean(np.equal(bd_label_p, bd_y_test))*100
print('Attack Success Rate for repaired_net_10:', asr_repaired_net)

```

Clean Classification accuracy for repaired\_net\_2: 96.04832424006236

Attack Success Rate for repaired\_net\_2: 100.0

-----

Clean Classification accuracy for repaired\_net\_4: 92.1278254091972

Attack Success Rate for repaired\_net\_4: 99.98441153546376

-----

Clean Classification accuracy for repaired\_net\_10: 84.3335931410756

Attack Success Rate for repaired\_net\_10: 77.20966484801247

[ ]: